

Lecture 10

May 26, 2023

1 Linear regression

1.1 Import necessary libraries

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

1.2 Import dataset

```
[2]: data = pd.read_csv("Salary_data.csv")
```

```
[3]: data.head()
```

```
[3]:  YearsExperience  Salary
0             1.1  39343.0
1             1.3  46205.0
2             1.5  37731.0
3             2.0  43525.0
4             2.2  39891.0
```

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes
```


```
[5]: data.isnull().any()
```

```
[5]: YearsExperience  False
Salary             False
```

dtype: bool

```
[6]: plt.scatter(data.YearsExperience, data.Salary)
plt.xlabel("Years")
plt.ylabel("Salary")
```

```
[6]: Text(0, 0.5, 'Salary')
```



1.3 Train test split

```
[17]: X = data.iloc[:, 0:1]
X
```

```
[17]:   YearsExperience
0          1.1
1          1.3
2          1.5
3          2.0
4          2.2
5          2.9
6          3.0
```

7	3.2
8	3.2
9	3.7
10	3.9
11	4.0
12	4.0
13	4.1
14	4.5
15	4.9
16	5.1
17	5.3
18	5.9
19	6.0
20	6.8
21	7.1
22	7.9
23	8.2
24	8.7
25	9.0
26	9.5
27	9.6
28	10.3
29	10.5

```
[18]: y = data.iloc[:, 1:]
```

```
[19]: y
```

```
[19]:      Salary
0    39343.0
1    46205.0
2    37731.0
3    43525.0
4    39891.0
5    56642.0
6    60150.0
7    54445.0
8    64445.0
9    57189.0
10   63218.0
11   55794.0
12   56957.0
13   57081.0
14   61111.0
15   67938.0
16   66029.0
17   83088.0
```

```
18 81363.0
19 93940.0
20 91738.0
21 98273.0
22 101302.0
23 113812.0
24 109431.0
25 105582.0
26 116969.0
27 112635.0
28 122391.0
29 121872.0
```

```
[20]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=0) #Splitting the dataset into train and test
```

```
[21]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(24, 1)
(6, 1)
(24, 1)
(6, 1)
```

1.4 Model building

```
[22]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
```

1.5 Train the model

```
[23]: lr.fit(X_train,y_train)
```

```
[23]: LinearRegression()
```

1.6 Test the model

```
[24]: y_pred=lr.predict(X_test)
```

```
[25]: y_pred #prediction
```

```
[25]: array([[ 40748.96184072],
[122699.62295594],
[ 64961.65717022],
[ 63099.14214487],
```

```
[115249.56285456],  
[107799.50275317]])
```

```
[26]: y_test # Actual outcome
```

```
[26]:      Salary  
2    37731.0  
28   122391.0  
13    57081.0  
10    63218.0  
26   116969.0  
24   109431.0
```

1.7 Error

```
[27]: E=y_test-y_pred  
E
```

```
[27]:      Salary  
2   -3017.961841  
28   -308.622956  
13  -7880.657170  
10   118.857855  
26   1719.437145  
24   1631.497247
```

1.8 R2 score

```
[28]: from sklearn.metrics import r2_score  
acc=r2_score(y_pred,y_test)  
acc
```

```
[28]: 0.9864826731176541
```

1.9 Predict random valuea

```
[29]: #Predict with random Values  
y_pr=lr.predict([[12]])  
y_pr
```

```
/Users/akashr/anaconda3/lib/python3.10/site-packages/sklearn/base.py:420:  
UserWarning: X does not have valid feature names, but LinearRegression was  
fitted with feature names  
warnings.warn(  


```

```
[29]: array([[138531.00067138]])
```

1.10 Plotting regression line


```
[30]: plt.scatter(X_train,y_train)  
plt.plot(X_train,lr.predict(X_train),'r')
```

[30]: [<matplotlib.lines.Line2D at 0x14d037cd0>]



```
[32]: plt.scatter(X_test,y_test)  
plt.plot(X_test,y_pred,'r')
```

[32]: [<matplotlib.lines.Line2D at 0x14d0ac250>]



2 Multilinear Regression

2.1 Import necessary libraries

```
[33]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

2.2 Import dataset

```
[34]: data = pd.read_csv("50_Startups.csv")
```

```
[35]: data.head()
```

```
[35]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
[36]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 108 entries, 0 to 107
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   R&D Spend             108 non-null   float64
1   Administration        108 non-null   float64
2   Marketing Spend       108 non-null   float64
3   State                 108 non-null   object
4   Profit                108 non-null   float64
dtypes: float64(4), object(1)
memory usage: 4.3+ KB
```

```
[37]: data.isnull().any()
```

```
[37]: R&D Spend           False
Administration        False
Marketing Spend       False
State                 False
Profit                False
dtype: bool
```

```
[38]: data.isnull().sum()
```

```
[38]: R&D Spend           0
Administration        0
Marketing Spend       0
State                 0
Profit                0
dtype: int64
```

```
[39]: data.State.value_counts()
```

```
[39]: New York           39
California            36
Florida              33
Name: State, dtype: int64
```

```
[40]: data.head()
```

```
[40]:   R&D Spend  Administration  Marketing Spend   State   Profit
0  165349.20      136897.80      471784.10  New York  192261.83
1  162597.70      151377.59      443898.53  California  191792.06
2  153441.51      101145.55      407934.54   Florida  191050.39
3  144372.41      118671.85      383199.62  New York  182901.99
4  142107.34       91391.77      366168.42   Florida  166187.94
```


2.3 Encoding

```
[41]: from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

```
[44]: data.State = le.fit_transform(data.State)
```

```
[45]: data.head()
```

```
[45]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	2	192261.83
1	162597.70	151377.59	443898.53	0	191792.06
2	153441.51	101145.55	407934.54	1	191050.39
3	144372.41	118671.85	383199.62	2	182901.99
4	142107.34	91391.77	366168.42	1	166187.94

```
[46]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 108 entries, 0 to 107  
Data columns (total 5 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   R&D Spend             108 non-null   float64  
1   Administration        108 non-null   float64  
2   Marketing Spend       108 non-null   float64  
3   State                 108 non-null   int64  
4   Profit                108 non-null   float64  
dtypes: float64(4), int64(1)  
memory usage: 4.3 KB
```

2.4 Dependent and Independent variables

```
[52]: y = data.loc[:, "Profit":]  
X = data.drop(columns=["Profit"], axis=1)
```

```
[53]: y
```

```
[53]:
```

	Profit
0	192261.83
1	191792.06
2	191050.39
3	182901.99
4	166187.94
..	...
103	132602.65
104	129917.04
105	126992.93

```
106 125370.37
107 124266.90
```

```
[108 rows x 1 columns]
```

```
[54]: X
```

```
[54]:
```

	R&D Spend	Administration	Marketing Spend	State
0	165349.20	136897.80	471784.10	2
1	162597.70	151377.59	443898.53	0
2	153441.51	101145.55	407934.54	1
3	144372.41	118671.85	383199.62	2
4	142107.34	91391.77	366168.42	1
..
103	119943.24	156547.42	256512.92	1
104	114523.61	122616.84	261776.23	2
105	78013.11	121597.55	264346.06	0
106	94657.16	145077.58	282574.31	2
107	91749.16	114175.79	294919.57	1

```
[108 rows x 4 columns]
```

2.5 Scaling

```
[50]: from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
```

```
[55]: columns = X.columns
```

```
[57]: X = scale.fit_transform(X)
X
```

```
[57]: array([[ 2.01819092,  0.55696553,  2.18783528,  1.16731536],
 [ 1.95675646,  1.0893969 ,  1.95029024, -1.23401909],
 [ 1.75232048, -0.75766814,  1.64392865, -0.03335187],
 [ 1.54982902, -0.1132146 ,  1.43322263,  1.16731536],
 [ 1.49925539, -1.11632102,  1.28814125, -0.03335187],
 [ 1.27083396, -0.806604 ,  1.25996984,  1.16731536],
 [ 1.3319795 ,  0.93574262, -0.74312422, -1.23401909],
 [ 1.23558377,  0.87437939,  0.92787631, -0.03335187],
 [ 1.01776417,  0.99163696,  0.82340982,  1.16731536],
 [ 1.08011094, -0.48065199,  0.76691751, -1.23401909],
 [ 0.60181303, -0.4102384 ,  0.1210342 , -0.03335187],
 [ 0.57410176, -1.10165542,  0.29637694, -1.23401909],
 [ 0.4220906 ,  0.20479752,  0.29718526, -0.03335187],
 [ 0.38030758,  0.50538623,  0.32125438, -1.23401909],
 [ 1.00438367,  1.27949493,  0.35403374, -0.03335187],
```

[0.88337621, 0.03184527, 0.39886959, 1.16731536],
 [0.06818346, -0.00563469, 0.42076085, -1.23401909],
 [0.43980557, 0.8577414 , 0.57603938, 1.16731536],
 [0.37487683, -0.27853767, 0.68120328, -0.03335187],
 [0.25588265, 1.16795814, -1.83108825, 1.16731536],
 [0.02890359, -0.28988105, 0.71310446, -1.23401909],
 [0.07658668, 1.1774935 , 0.72224334, 1.16731536],
 [-0.02154122, 0.03794589, 0.75275659, -0.03335187],
 [-0.165823 , -0.58832162, 0.76510399, -0.03335187],
 [0.04654576, -0.82621636, -0.63359259, 1.16731536],
 [-0.22985461, 0.65460485, -0.6558447 , -1.23401909],
 [0.00825076, 0.82311815, -0.68917401, -0.03335187],
 [-0.06367255, 0.224807 , 1.177531 , 1.16731536],
 [-0.19889043, 2.23914065, -0.82463513, -0.03335187],
 [-0.20884945, 1.15023284, -0.918423 , 1.16731536],
 [-0.2894745 , -0.22465065, -1.05478086, -0.03335187],
 [-0.30863384, 1.13809338, -1.07959552, 1.16731536],
 [-0.25789476, 0.27463345, -1.43850805, -1.23401909],
 [-0.4346159 , -0.68736485, -0.00270767, -0.03335187],
 [-0.63708012, 1.32165249, -0.03539459, -1.23401909],
 [-0.64628022, -1.3496062 , -0.08037287, 1.16731536],
 [-1.03367032, 0.19508381, -0.11777636, -0.03335187],
 [-0.68968668, -2.59114164, -0.15268033, -1.23401909],
 [-1.2219853 , -2.0519077 , -0.25289545, 1.16731536],
 [-0.81274405, -1.42555047, -0.34034534, -1.23401909],
 [-1.0316481 , -0.11784035, -0.3591171 , -1.23401909],
 [-1.05088134, -1.36198577, -0.43003381, -0.03335187],
 [-1.14581818, -0.93990057, -0.57033117, -1.23401909],
 [-1.32745787, 0.20707435, -1.52838812, 1.16731536],
 [-1.17848774, 1.21546693, -1.5897171 , -1.23401909],
 [-1.65133124, 0.08833235, -1.8148695 , 1.16731536],
 [-1.6442929 , -0.21821836, 0.6999006 , -0.03335187],
 [-1.673664 , 0.50288031, -1.83108825, -1.23401909],
 [-1.66156131, -2.57422677, -1.83108825, 1.16731536],
 [-1.673664 , -0.17528531, -1.4462786 , -1.23401909],
 [0.57410176, -1.10165542, 0.29637694, -1.23401909],
 [0.4220906 , 0.20479752, 0.29718526, -0.03335187],
 [0.38030758, 0.50538623, 0.32125438, -1.23401909],
 [1.00438367, 1.27949493, 0.35403374, -0.03335187],
 [0.88337621, 0.03184527, 0.39886959, 1.16731536],
 [0.06818346, -0.00563469, 0.42076085, -1.23401909],
 [0.43980557, 0.8577414 , 0.57603938, 1.16731536],
 [0.37487683, -0.27853767, 0.68120328, -0.03335187],
 [0.25588265, 1.16795814, -1.83108825, 1.16731536],
 [0.02890359, -0.28988105, 0.71310446, -1.23401909],
 [0.07658668, 1.1774935 , 0.72224334, 1.16731536],
 [-0.02154122, 0.03794589, 0.75275659, -0.03335187],

```

[-0.165823 , -0.58832162, 0.76510399, -0.03335187],
[ 0.04654576, -0.82621636, -0.63359259, 1.16731536],
[-0.22985461, 0.65460485, -0.6558447 , -1.23401909],
[ 0.00825076, 0.82311815, -0.68917401, -0.03335187],
[-0.06367255, 0.224807 , 1.177531 , 1.16731536],
[-0.19889043, 2.23914065, -0.82463513, -0.03335187],
[-0.20884945, 1.15023284, -0.918423 , 1.16731536],
[-0.2894745 , -0.22465065, -1.05478086, 1.16731536],
[-1.2219853 , 1.13809338, -1.07959552, 1.16731536],
[-0.25789476, 0.27463345, -1.43850805, -1.23401909],
[-0.4346159 , -0.68736485, -0.00270767, -0.03335187],
[-0.63708012, 1.32165249, -0.03539459, 1.16731536],
[-1.2219853 , -1.3496062 , -0.08037287, 1.16731536],
[-1.03367032, 0.19508381, -0.11777636, -0.03335187],
[-0.68968668, -2.59114164, -0.15268033, -1.23401909],
[-1.2219853 , -2.0519077 , -0.25289545, 1.16731536],
[-0.81274405, -1.42555047, -0.34034534, -1.23401909],
[-1.0316481 , -0.11784035, -0.3591171 , -1.23401909],
[-1.05088134, -1.36198577, -0.43003381, 1.16731536],
[-1.14581818, -0.93990057, -0.57033117, -1.23401909],
[-1.2219853 , 0.20707435, -1.52838812, 1.16731536],
[-1.17848774, 1.21546693, -1.5897171 , -1.23401909],
[-1.65133124, 0.08833235, -1.8148695 , 1.16731536],
[-1.6442929 , -0.21821836, 0.6999006 , -0.03335187],
[-1.673664 , 0.50288031, -1.83108825, -1.23401909],
[-1.66156131, -2.57422677, -1.83108825, 1.16731536],
[-1.673664 , -0.17528531, -1.4462786 , -1.23401909],
[ 2.01819092, 0.55696553, 2.18783528, 1.16731536],
[ 1.95675646, 1.0893969 , 1.95029024, -1.23401909],
[ 1.75232048, -0.75766814, 1.64392865, -0.03335187],
[ 1.54982902, -0.1132146 , 1.43322263, 1.16731536],
[ 1.49925539, -1.11632102, 1.28814125, -0.03335187],
[ 1.27083396, -0.806604 , 1.25996984, 1.16731536],
[ 1.3319795 , 0.93574262, -0.74312422, -1.23401909],
[ 1.23558377, 0.87437939, 0.92787631, -0.03335187],
[ 1.01776417, 0.99163696, 0.82340982, 1.16731536],
[ 1.08011094, -0.48065199, 0.76691751, -1.23401909],
[ 0.60181303, -0.4102384 , 0.1210342 , -0.03335187],
[ 0.57410176, -1.10165542, 0.29637694, -1.23401909],
[ 0.4220906 , 0.20479752, 0.29718526, -0.03335187],
[ 0.38030758, 0.50538623, 0.32125438, -1.23401909],
[ 1.00438367, 1.27949493, 0.35403374, -0.03335187],
[ 0.88337621, 0.03184527, 0.39886959, 1.16731536],
[ 0.06818346, -0.00563469, 0.42076085, -1.23401909],
[ 0.43980557, 0.8577414 , 0.57603938, 1.16731536],
[ 0.37487683, -0.27853767, 0.68120328, -0.03335187]]

```

```
[58]: X= pd.DataFrame(X, columns=columns)
X
```

```
[58]:
```

	R&D Spend	Administration	Marketing Spend	State
0	2.018191	0.556966	2.187835	1.167315
1	1.956756	1.089397	1.950290	-1.234019
2	1.752320	-0.757668	1.643929	-0.033352
3	1.549829	-0.113215	1.433223	1.167315
4	1.499255	-1.116321	1.288141	-0.033352
..
103	1.004384	1.279495	0.354034	-0.033352
104	0.883376	0.031845	0.398870	1.167315
105	0.068183	-0.005635	0.420761	-1.234019
106	0.439806	0.857741	0.576039	1.167315
107	0.374877	-0.278538	0.681203	-0.033352

[108 rows x 4 columns]

2.6 Train test split

```
[59]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=0)#Splitting the dataset into train and test
```

```
[60]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(86, 4)
(22, 4)
(86, 1)
(22, 1)
```

2.7 Model building

```
[61]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
```

2.8 Train the model

```
[62]: lr.fit(X_train,y_train)
```

```
[62]: LinearRegression()
```

2.9 Test the model

```
[63]: y_pred=lr.predict(X_test)
```

```
[64]: y_pred #prediction
```

```
[64]: array([[ 48379.24868386],  
            [134848.9192467 ],  
            [ 76483.10965225],  
            [181561.78529188],  
            [112966.00035112],  
            [134241.68257226],  
            [129222.38731942],  
            [160017.05919784],  
            [116752.96050775],  
            [ 46260.9926407 ],  
            [102266.36060189],  
            [115568.27177598],  
            [ 48379.24868386],  
            [119114.16525526],  
            [ 88602.22002789],  
            [127101.40273623],  
            [127101.40273623],  
            [ 90942.33421515],  
            [ 58677.7334991 ],  
            [146304.26215166],  
            [149410.05441   ],  
            [152502.87810448]])
```

```
[65]: y_test # Actual outcome
```

```
[65]:      Profit  
84      64926.08  
10     146121.95  
75      90708.19  
2      191050.39  
24     108552.04  
100    144259.40  
107    124266.90  
7      155752.60  
16     126992.93  
86      42559.73  
68     101004.64  
22     110352.25  
45      64926.08  
60     111313.02  
76      89949.14  
52     134307.35
```

```
13    134307.35
73     96712.80
85     49490.75
54    129917.04
103   132602.65
8     152211.77
```

2.10 Error

```
[66]: E=y_test-y_pred
      E
```

```
[66]:          Profit
84    16546.831316
10    11273.030753
75    14225.080348
2      9488.604708
24   -4413.960351
100   10017.717428
107   -4955.487319
7     -4264.459198
16    10239.969492
86    -3701.262641
68    -1261.720602
22    -5216.021776
45    16546.831316
60    -7801.145255
76     1346.919972
52     7205.947264
13     7205.947264
73     5770.465785
85    -9186.983499
54   -16387.222152
103  -16807.404410
8     -291.108104
```

3 Polynomial regression

3.1 Import necessary libraries

```
[1]: import numpy as np
      import pandas as pd
      import seaborn as sns
      import matplotlib.pyplot as plt
```

3.2 Import dataset

```
[68]: data = pd.read_csv("Position_Salaries.csv")
```

```
[69]: data.head()
```

```
[69]:
```

	Position	Level	Salary
0	Business Analyst	1	45000
1	Junior Consultant	2	50000
2	Senior Consultant	3	60000
3	Manager	4	80000
4	Country Manager	5	110000

```
[70]: data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Position    10 non-null     object
1   Level       10 non-null     int64
2   Salary      10 non-null     int64
dtypes: int64(2), object(1)
memory usage: 368.0+ bytes
```

```
[71]: data.isnull().any()
```

```
[71]: Position    False
Level         False
Salary        False
dtype: bool
```

```
[73]: plt.scatter(data.Level, data.Salary)
plt.xlabel("Level")
plt.ylabel("Salary")
```

```
[73]: Text(0, 0.5, 'Salary')
```

3.3 Train test split

```
[74]: X = data.iloc[:, 1:2]  
X
```

```
[74]:   Level  
0      1  
1      2  
2      3  
3      4  
4      5  
5      6  
6      7  
7      8  
8      9  
9     10
```

```
[75]: y = data.iloc[:, 2:3]
```

```
[76]: y
```

```
[76]:      Salary
      0    45000
      1    50000
      2    60000
      3    80000
      4   110000
      5   150000
      6   200000
      7   300000
      8   500000
      9  1000000
```

```
[77]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=0) #Splitting the dataset into train and test
```

```
[78]: print(X_train.shape)
      print(X_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(8, 1)
(2, 1)
(8, 1)
(2, 1)
```

3.4 Model building

```
[79]: from sklearn.linear_model import LinearRegression
      lr=LinearRegression()
```

3.5 Train the model

```
[80]: lr.fit(X_train,y_train)
```

```
[80]: LinearRegression()
```

3.6 Test the model

```
[81]: y_pred=lr.predict(X_test)
```

```
[82]: y_pred #prediction
```

```
[82]: array([[ 44275.9295499 ],
      [543473.58121331]])
```

```
[83]: y_test # Actual outcome
```

```
[83]:      Salary
      2    60000
      8   500000
```

3.7 Error

```
[84]: E=y_test-y_pred
      E
```

```
[84]:      Salary
      2  15724.070450
      8 -43473.581213
```

3.8 R2 score

```
[85]: from sklearn.metrics import r2_score
      acc=r2_score(y_pred,y_test)
      acc
```

```
[85]: 0.982847405505933
```

3.9 Predict random valuea

```
[86]: #Predict with random Values
      y_pr=lr.predict([[12]])
      y_pr
```


```
/Users/akashr/anaconda3/lib/python3.10/site-packages/sklearn/base.py:420:
UserWarning: X does not have valid feature names, but LinearRegression was
fitted with feature names
  warnings.warn(
```

```
[86]: array([[793072.40704501]])
```

3.10 Plotting regression line


```
[87]: plt.scatter(X_train,y_train)
      plt.plot(X_train,lr.predict(X_train),'r')
```

```
[87]: [<matplotlib.lines.Line2D at 0x14f3d1270>]
```



```
[88]: plt.scatter(X_test,y_test)
      plt.plot(X_test,y_pred,'r')
```

```
[88]: [<matplotlib.lines.Line2D at 0x14f43d930>]
```



3.11 R2 score

```
[89]: from sklearn.metrics import r2_score
      acc=r2_score(y_pred,y_test)
      acc
```

[89]: 0.982847405505933

3.12 Adding Polynomial features

```
[90]: #convert the normal values to polynomial features
      from sklearn.preprocessing import PolynomialFeatures
      pr=PolynomialFeatures(degree=4)
      #x0,x1,x2,x3,x4
      X_poly=pr.fit_transform(X)
      X_poly
```

```
[90]: array([[1.000e+00, 1.000e+00, 1.000e+00, 1.000e+00, 1.000e+00],
             [1.000e+00, 2.000e+00, 4.000e+00, 8.000e+00, 1.600e+01],
             [1.000e+00, 3.000e+00, 9.000e+00, 2.700e+01, 8.100e+01],
             [1.000e+00, 4.000e+00, 1.600e+01, 6.400e+01, 2.560e+02],
             [1.000e+00, 5.000e+00, 2.500e+01, 1.250e+02, 6.250e+02],
```

```
[1.000e+00, 6.000e+00, 3.600e+01, 2.160e+02, 1.296e+03],
[1.000e+00, 7.000e+00, 4.900e+01, 3.430e+02, 2.401e+03],
[1.000e+00, 8.000e+00, 6.400e+01, 5.120e+02, 4.096e+03],
[1.000e+00, 9.000e+00, 8.100e+01, 7.290e+02, 6.561e+03],
[1.000e+00, 1.000e+01, 1.000e+02, 1.000e+03, 1.000e+04]])
```

```
[91]: #associating above values with y
pr.fit(X_poly,y)
```

```
[91]: PolynomialFeatures(degree=4)
```

```
[93]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()

lr.fit(X_poly,y)

pred=lr.predict(X_poly)
```

```
[94]: pred
```

```
[94]: array([[ 53356.6433568 ],
[ 31759.90675996],
[ 58642.19114221],
[ 94632.86713289],
[121724.94172497],
[143275.05827508],
[184003.49650349],
[289994.17249411],
[528694.63869453],
[988916.08391595]])
```

```
[95]: y
```

```
[95]:      Salary
0      45000
1      50000
2      60000
3      80000
4     110000
5     150000
6     200000
7     300000
8     500000
9    1000000
```

```
[96]: plt.scatter(X,y)
plt.plot(X,pred,'r')
```

[96]: [

