

Lecture 15

June 2, 2023

```
[1]: import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
```

```
[2]: sns.get_dataset_names()
```

```
[2]: ['anagrams',
      'anscombe',
      'attention',
      'brain_networks',
      'car_crashes',
      'diamonds',
      'dots',
      'dowjones',
      'exercise',
      'flights',
      'fmri',
      'geyser',
      'glue',
      'healthexp',
      'iris',
      'mpg',
      'penguins',
      'planets',
      'seaice',
      'taxis',
      'tips',
      'titanic']
```

```
[3]: df=sns.load_dataset('tips')
```

```
[4]: df.head()
```

```
[4]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3

3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

```
[5]: df.isnull().sum()
```

```
[5]: total_bill    0
      tip          0
      sex          0
      smoker       0
      day          0
      time         0
      size         0
      dtype: int64
```

```
[6]: from sklearn.preprocessing import LabelEncoder
      le=LabelEncoder()
```

```
[7]: df['sex']=le.fit_transform(df['sex'])
      df['smoker']=le.fit_transform(df['smoker'])
      df['day']=le.fit_transform(df['day'])
      df['time']=le.fit_transform(df['time'])
```

```
[8]: df.head()
```

```
[8]:   total_bill  tip  sex  smoker  day  time  size
0      16.99  1.01   0      0     2     0     2
1      10.34  1.66   1      0     2     0     3
2      21.01  3.50   1      0     2     0     3
3      23.68  3.31   1      0     2     0     2
4      24.59  3.61   0      0     2     0     4
```

1 Splitting dataset into X and Y

```
[9]: x=df.drop('tip',axis=1)
      y=df['tip']
```

```
[10]: x
```

```
[10]:   total_bill  sex  smoker  day  time  size
0      16.99   0      0     2     0     2
1      10.34   1      0     2     0     3
2      21.01   1      0     2     0     3
3      23.68   1      0     2     0     2
4      24.59   0      0     2     0     4
..      ...   ...      ...   ...   ...   ...
239     29.03   1      0     1     0     3
240     27.18   0      1     1     0     2
```

```

241      22.67      1      1      1      0      2
242      17.82      1      0      1      0      2
243      18.78      0      0      3      0      2

```

[244 rows x 6 columns]

```
[11]: y
```

```

[11]: 0      1.01
      1      1.66
      2      3.50
      3      3.31
      4      3.61

```

```

...
239      5.92
240      2.00
241      2.00
242      1.75
243      3.00

```

Name: tip, Length: 244, dtype: float64

2 Applying Scaling on X

```
[12]: from sklearn.preprocessing import MinMaxScaler
```

```
[13]: scale=MinMaxScaler()
```

```
[14]: x=scale.fit_transform(x)
```

```
[15]: x
```

```

[15]: array([[0.29157939, 0.        , 0.        , 0.66666667, 0.        ,
              0.2        ],
             [0.1522832 , 1.        , 0.        , 0.66666667, 0.        ,
              0.4        ],
             [0.3757855 , 1.        , 0.        , 0.66666667, 0.        ,
              0.4        ],
             ...,
             [0.41055718, 1.        , 1.        , 0.33333333, 0.        ,
              0.2        ],
             [0.30896523, 1.        , 0.        , 0.33333333, 0.        ,
              0.2        ],
             [0.32907415, 0.        , 0.        , 1.        , 0.        ,
              0.2        ]])

```

3 Splitting X and Y in train and Test

```
[16]: from sklearn.model_selection import train_test_split
```

```
[17]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

4 Model Building

```
[18]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()
```

```
[19]: lr.fit(x_train,y_train)
```

```
[19]: LinearRegression()
```

```
[20]: y_pred=lr.predict(x_test)
```

5 Evaluation metrics

```
[21]: from sklearn.metrics import mean_squared_error,r2_score, mean_absolute_error
```

```
[22]: error=y_test-y_pred
```

```
[23]: error
```

```
[23]: 64      -0.268407  
63       0.694602  
55       0.762951  
111     -0.476576  
225     -0.238789  
...  
90      -0.668911  
101      0.591909  
75      -0.783550  
4       -0.036714  
109      1.789722  
Name: tip, Length: 74, dtype: float64
```

```
[24]: se=error*error
```

```
[25]: se
```

```
[25]: 64      0.072042  
63      0.482472  
55      0.582094  
111     0.227124
```

```

225    0.057020
...
90     0.447442
101    0.350356
75     0.613951
4      0.001348
109    3.203105
Name: tip, Length: 74, dtype: float64

```

```
[26]: mse=np.mean(se)
```

```
[27]: mse
```

```
[27]: 0.9166571859645408
```

```
[28]: mse2=mean_squared_error(y_test,y_pred)
```

```
[29]: mse2
```

```
[29]: 0.9166571859645408
```

```
[30]: mae=mean_absolute_error(y_test,y_pred)
```

```
[31]: mae
```

```
[31]: 0.7166503728723608
```

```
[32]: rmse=np.sqrt(mse2)
```

```
[33]: rmse
```

```
[33]: 0.9574221566083275
```

```
[34]: r2=r2_score(y_test,y_pred)
```

```
[35]: r2
```

```
[35]: 0.4687997604863058
```

6 Classification Metrics

```
[37]: data=sns.load_dataset('titanic')
```

```
[38]: data
```

```
[38]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	\
0	0	3	male	22.0	1	0	7.2500	S	Third	
1	1	1	female	38.0	1	0	71.2833	C	First	

2	1	3	female	26.0	0	0	7.9250	S	Third
3	1	1	female	35.0	1	0	53.1000	S	First
4	0	3	male	35.0	0	0	8.0500	S	Third
..
886	0	2	male	27.0	0	0	13.0000	S	Second
887	1	1	female	19.0	0	0	30.0000	S	First
888	0	3	female	NaN	1	2	23.4500	S	Third
889	1	1	male	26.0	0	0	30.0000	C	First
890	0	3	male	32.0	0	0	7.7500	Q	Third

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True
..
886	man	True	NaN	Southampton	no	True
887	woman	False	B	Southampton	yes	True
888	woman	False	NaN	Southampton	no	False
889	man	True	C	Cherbourg	yes	True
890	man	True	NaN	Queenstown	no	True

[891 rows x 15 columns]

```
[39]: data.isnull().sum()
```

```
[39]: survived      0
pclass            0
sex              0
age             177
sibsp           0
parch           0
fare            0
embarked        2
class           0
who             0
adult_male      0
deck           688
embark_town     2
alive           0
alone           0
dtype: int64
```

```
[40]: data['age'].fillna(data['age'].median(),inplace=True)
```

```
[41]: data['deck'].mode()
```

```
[41]: 0    C
      Name: deck, dtype: category
      Categories (7, object): ['A', 'B', 'C', 'D', 'E', 'F', 'G']
```

```
[42]: data['deck'].fillna(data['deck'].mode()[0],inplace=True)
```

```
[43]: data.isnull().sum()
```

```
[43]: survived      0
      pclass       0
      sex          0
      age          0
      sibsp        0
      parch        0
      fare         0
      embarked     2
      class        0
      who          0
      adult_male   0
      deck         0
      embark_town  2
      alive        0
      alone        0
      dtype: int64
```

```
[44]: data.dropna(how='any',inplace=True)
```

```
[45]: data.shape
```

```
[45]: (889, 15)
```

```
[46]: data.isnull().sum()
```

```
[46]: survived      0
      pclass       0
      sex          0
      age          0
      sibsp        0
      parch        0
      fare         0
      embarked     0
      class        0
      who          0
      adult_male   0
      deck         0
      embark_town  0
      alive        0
      alone        0
```

dtype: int64

```
[47]: data.head()
```

```
[47]:   survived  pclass    sex  age  sibsp  parch    fare embarked  class \
0         0       3   male  22.0     1     0   7.2500         S   Third
1         1       1  female  38.0     1     0  71.2833         C   First
2         1       3  female  26.0     0     0   7.9250         S   Third
3         1       1  female  35.0     1     0  53.1000         S   First
4         0       3   male  35.0     0     0   8.0500         S   Third

      who  adult_male deck  embark_town  alive  alone
0    man         True   C  Southampton    no  False
1  woman        False   C   Cherbourg   yes  False
2  woman        False   C  Southampton   yes   True
3  woman        False   C  Southampton   yes  False
4    man         True   C  Southampton    no   True
```

```
[48]: data.drop('embark_town',axis=1,inplace=True)
```

```
[49]: data['sex']=le.fit_transform(data['sex'])
data['embarked']=le.fit_transform(data['embarked'])
data['class']=le.fit_transform(data['class'])
data['who']=le.fit_transform(data['who'])
data['adult_male']=le.fit_transform(data['adult_male'])
data['deck']=le.fit_transform(data['deck'])
data['alive']=le.fit_transform(data['alive'])
data['alone']=le.fit_transform(data['alone'])
```

```
[50]: data.head()
```

```
[50]:   survived  pclass  sex  age  sibsp  parch    fare  embarked  class  who \
0         0       3    1  22.0     1     0   7.2500         2      2    1
1         1       1    0  38.0     1     0  71.2833         0      0    2
2         1       3    0  26.0     0     0   7.9250         2      2    2
3         1       1    0  35.0     1     0  53.1000         2      0    2
4         0       3    1  35.0     0     0   8.0500         2      2    1

      adult_male  deck  alive  alone
0             1     2      0      0
1             0     2      1      0
2             0     2      1      1
3             0     2      1      0
4             1     2      0      1
```

```
[51]: x=data.drop('survived',axis=1)
y=data['survived']
```



```
[52]: x
```

```
[52]:      pclass  sex  age  sibsp  parch    fare  embarked  class  who  \
0         3    1  22.0     1     0   7.2500         2     2    1
1         1    0  38.0     1     0  71.2833         0     0    2
2         3    0  26.0     0     0   7.9250         2     2    2
3         1    0  35.0     1     0  53.1000         2     0    2
4         3    1  35.0     0     0   8.0500         2     2    1
..      ...  ...  ...  ...  ...  ...  ...  ...
886        2    1  27.0     0     0  13.0000         2     1    1
887        1    0  19.0     0     0  30.0000         2     0    2
888        3    0  28.0     1     2  23.4500         2     2    2
889        1    1  26.0     0     0  30.0000         0     0    1
890        3    1  32.0     0     0   7.7500         1     2    1

      adult_male  deck  alive  alone
0              1     2     0     0
1              0     2     1     0
2              0     2     1     1
3              0     2     1     0
4              1     2     0     1
..      ...  ...  ...  ...
886          1     2     0     1
887          0     1     1     1
888          0     2     0     0
889          1     2     1     1
890          1     2     0     1
```

[889 rows x 13 columns]

```
[53]: y
```

```
[53]: 0      0
1      1
2      1
3      1
4      0
..
886    0
887    1
888    0
889    1
890    0
Name: survived, Length: 889, dtype: int64
```

```
[54]: x=scale.fit_transform(x)
```

```
[55]: x
```

```
[55]: array([[1.      , 1.      , 0.27117366, ..., 0.33333333, 0.      ,
          0.      ],
          [0.      , 0.      , 0.4722292 , ..., 0.33333333, 1.      ,
          0.      ],
          [1.      , 0.      , 0.32143755, ..., 0.33333333, 1.      ,
          1.      ],
          ...,
          [1.      , 0.      , 0.34656949, ..., 0.33333333, 0.      ,
          0.      ],
          [0.      , 1.      , 0.32143755, ..., 0.33333333, 1.      ,
          1.      ],
          [1.      , 1.      , 0.39683338, ..., 0.33333333, 0.      ,
          1.      ]])
```

```
[56]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

```
[57]: from sklearn.linear_model import LogisticRegression
      lor=LogisticRegression()
```

```
[58]: lor.fit(x_train,y_train)
```

```
[58]: LogisticRegression()
```

```
[59]: y_pred=lor.predict(x_test)
```

```
[60]: from sklearn.metrics import
      ↪accuracy_score,confusion_matrix,roc_auc_score,precision_recall_fscore_support,f1_score,roc_
```

```
[61]: accuracy_score(y_test,y_pred)
```

```
[61]: 1.0
```

```
[62]: confusion_matrix(y_test,y_pred)
```

```
[62]: array([[157,  0],
          [ 0, 110]])
```

```
[63]: pr=precision_recall_fscore_support(y_test,y_pred, average='micro')
```

```
[64]: pr
```

```
[64]: (1.0, 1.0, 1.0, None)
```

```
[65]: f1_score(y_test,y_pred)
```

```
[65]: 1.0
```

```
[66]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	157
1	1.00	1.00	1.00	110
accuracy			1.00	267
macro avg	1.00	1.00	1.00	267
weighted avg	1.00	1.00	1.00	267


```
[67]: roc=roc_auc_score(y_test,y_pred)
```

```
[68]: roc
```

```
[68]: 1.0
```

```
[69]: fpr,tpr,threshold=roc_curve(y_test,y_pred)
plt.plot(fpr,tpr,marker='.')
plt.xlabel('False Positive Rate')
plt.ylabel('Ture Positive Rate')
```

```
[69]: Text(0, 0.5, 'Ture Positive Rate')
```



[]: